

DTSA5509 Supervised Machine Learning Final Project

Location of this project: <https://github.com/NikoKuu/Niko-s-intro-to-ML-project>

Project Topic

A dataset about heart attack risk was somewhat arbitrarily chosen for this project. *The objective of this project is to create a model that can predict, when given a set of inputs, whether a person is in a high risk category of experiencing a heart attack.* Having an accurate model for this data is literally a matter of life and death as certain interventions could be done for high-risk patients.

We will be trying to model the data with multiple supervised machine learning algorithms covered in the course:

1. 'Traditional' Logistic Regression,
2. (Pruned) Decision Tree,
3. Ensemble method AdaBoost and
4. Support Vector Machine model.

If a method has hyperparameters or other tuning parameters, best effort is made to get the best result. Then, based on the calculated metrics, the best model is chosen.

Data

The data set was found on Kaggle. The full source of the data set:

Rashik Rahman. (2021?). Heart Attack Analysis & Prediction Dataset, Version 2.

Retrieved 09/06/2024 from <https://www.kaggle.com/datasets/rashikrahmanpritom/heart-attack-analysis-prediction-dataset/version/2>.

The data set consists of the occurrence of a heart attack and 13 associated risk factors. The plan is to train a model using the risk factors as labels and whether heart attack happened or not as a target. That makes this work a binary classifier model. The thirteen features fall in three categories: binary (3), categorical (5) and continuous (5). The raw data has 303 samples (rows). The data set was saved in csv format and Pandas Python package was able to read it without any problems.

Data Cleaning

In the code below, basic information about the dataset is printed out: Size of the dataset, data types of the variables. Also, a couple of first rows of the dataset is printed out to show how the dataset looks.

Descriptive statistics show if there're unreasonable values (min and max values) that are probably invalid data entries. It also shows means and standard deviations.

Another check for invalid, not-a-number, values shows that the data set is clean. However, when looking for duplicate entries, one row was found to be a duplicate and thus removed.

Finally, we checked how evenly the output variable is distributed. The response/target variable is distributed evenly between had not heart attack and had heart attack.

Later, in the EDA section, it was found that the Old Peak measurement ('oldpeak') had too many zero values to be valid and those values were imputed with mean value of the feature. Another option would have been to drop the feature

```
In [425... import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

datafile = 'heart.csv'
df = pd.read_csv(datafile)

print('\n*** Dataframe info ***')
print(df.info())
print('\n*** Five first rows ***')
print(df.head(5))
#print(df.columns[0])
#print('\nDatatypes in the dataframe:')
#print(df.dtypes)
print('\n*** Descriptive statistics ***')
print(df.describe())

#for c in df.columns[0:]:
#    print(c, df[c].unique())

#df['age'] = pd.to_numeric(df['age'], errors='coerce')
print('\n*** Any NaN values ***')
print(df.isna().any())

print('\n*** Any duplicate rows? If so, remove them. ***')
print('Duplicate rows found:')
print(df[df.duplicated(keep='last')])
before = len(df)
#print(df.loc[163:164])
df = df.drop_duplicates()
print('Number of rows before', before, 'and after', len(df), 'removing duplicates.')

print('\n*** Balance of the response variable ***')
print('Number of more chance of heart attack:', sum(df.output==1), '('+str(round(sum(df.output==1)/len(df)*10
print('Number of less chance of heart attack:', sum(df.output==0), '('+str(round(sum(df.output==0)/len(df)*10

#print(pd.unique(df['slp']))
```

```

*** Dataframe info ***
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null   int64
1   sex         303 non-null   int64
2   cp          303 non-null   int64
3   trtbps      303 non-null   int64
4   chol        303 non-null   int64
5   fbs         303 non-null   int64
6   restecg     303 non-null   int64
7   thalachh    303 non-null   int64
8   exng        303 non-null   int64
9   oldpeak     303 non-null   float64
10  slp         303 non-null   int64
11  caa         303 non-null   int64
12  thall       303 non-null   int64
13  output      303 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
None

```

```

*** Five first rows ***
   age  sex  cp  trtbps  chol  fbs  restecg  thalachh  exng  oldpeak  slp  \
0   63   1   3    145   233    1         0     150     0     2.3     0
1   37   1   2    130   250    0         1     187     0     3.5     0
2   41   0   1    130   204    0         0     172     0     1.4     2
3   56   1   1    120   236    0         1     178     0     0.8     2
4   57   0   0    120   354    0         1     163     1     0.6     2

```

```

   caa  thall  output
0     0      1       1
1     0      2       1
2     0      2       1
3     0      2       1
4     0      2       1

```

```

*** Descriptive statistics ***
          age          sex          cp          trtbps          chol          fbs  \
count  303.000000  303.000000  303.000000  303.000000  303.000000  303.000000
mean    54.366337   0.683168   0.966997  131.623762  246.264026   0.148515
std     9.082101   0.466011   1.032052  17.538143   51.830751   0.356198
min     29.000000   0.000000   0.000000   94.000000  126.000000   0.000000
25%     47.500000   0.000000   0.000000  120.000000  211.000000   0.000000
50%     55.000000   1.000000   1.000000  130.000000  240.000000   0.000000
75%     61.000000   1.000000   2.000000  140.000000  274.500000   0.000000
max     77.000000   1.000000   3.000000  200.000000  564.000000   1.000000

```

```

          restecg  thalachh          exng  oldpeak          slp          caa  \
count  303.000000  303.000000  303.000000  303.000000  303.000000  303.000000
mean     0.528053  149.646865   0.326733   1.039604   1.399340   0.729373
std     0.525860   22.905161   0.469794   1.161075   0.616226   1.022606
min      0.000000   71.000000   0.000000   0.000000   0.000000   0.000000
25%      0.000000  133.500000   0.000000   0.000000   1.000000   0.000000
50%      1.000000  153.000000   0.000000   0.800000   1.000000   0.000000
75%      1.000000  166.000000   1.000000   1.600000   2.000000   1.000000
max      2.000000  202.000000   1.000000   6.200000   2.000000   4.000000

```

```

          thall  output
count  303.000000  303.000000
mean     2.313531   0.544554
std     0.612277   0.498835
min      0.000000   0.000000
25%      2.000000   0.000000

```

```
50%      2.000000    1.000000
75%      3.000000    1.000000
max       3.000000    1.000000
```

*** Any NaN values ***

```
age      False
sex      False
cp       False
trtbps   False
chol     False
fbs      False
restecg  False
thalachh False
exng     False
oldpeak  False
slp      False
caa      False
thall    False
output   False
dtype: bool
```

*** Any duplicate rows? If so, remove them. ***

Duplicate rows found:

```
   age sex  cp  trtbps  chol  fbs  restecg  thalachh  exng  oldpeak  slp  \
163   38   1   2    138   175    0         1       173    0      0.0    2
      caa  thall  output
163    4     2        1
```

Number of rows before 303 and after 302 removing duplicates.

*** Balance of the response variable ***

Number of more chance of heart attack: 164 (54%)

Number of less chance of heart attack: 138 (46%)

Dataset info

The feature explanations (as stated on the Kaggle page) and datatypes are found below:

0. Age: Age of the patient
 - Real-valued
1. Sex: Gender of the person
 - Binary-valued
2. cp: Chest Pain type
 - Categorical values
 - Value 1: typical angina
 - Value 2: atypical angina
 - Value 3: non-anginal pain
 - Value 4: asymptomatic
3. trtbps: resting blood pressure (in mm Hg)
 - Real-valued
4. chol: cholestoral in mg/dl fetched via BMI sensor
 - Real-valued
5. fbs: (fasting blood sugar > 120 mg/dl)
 - Binary-valued
 - Value 1: true
 - Value 0: false
6. rest_ecg: resting electrocardiographic results
 - Categorical values
 - Value 0: normal

- Value 1: having ST-T wave abnormality (T wave inversions and/or ST elevation or depression of > 0.05 mV)
 - Value 2: showing probable or definite left ventricular hypertrophy by Estes' criteria
7. thalach: maximum heart rate achieved
- Real-valued
8. exang: exercise induced angina
- Binary-valued
 - Value 1: Yes
 - Value 0: No
9. oldpeak: Previous peak
- Real-valued
10. slp: Slope
- Categorical values: 0, 1, 2
11. ca: number of major vessels (0-3)
- Real-valued integer
12. thall: Thal rate
- Categorical values: 0, 1, 2, 3
13. target:
- Binary-valued
 - Value 0: less chance of heart attack
 - Value 1: more chance of heart attack

Exploratory Data Analysis

In EDA, we will change the data types of the binary and categorical features. Then, we check each the distribution and potential erroneous or missing values of each feature. Starting from the binary-valued to categorical and finally real-valued features.

In the code below, we change the 'int64' data types to 'category' data types. This does not change the modeling process but makes it easier to identify what kind of data is in each feature. For convenience, the original dataframe is also preserved.

```
In [426... df_cat = df.copy()
df_cat = df_cat.astype({'sex':'category', 'cp':'category', 'fbs':'category', 'restecg':'category', 'exng':'ca
df_cat.info()
```

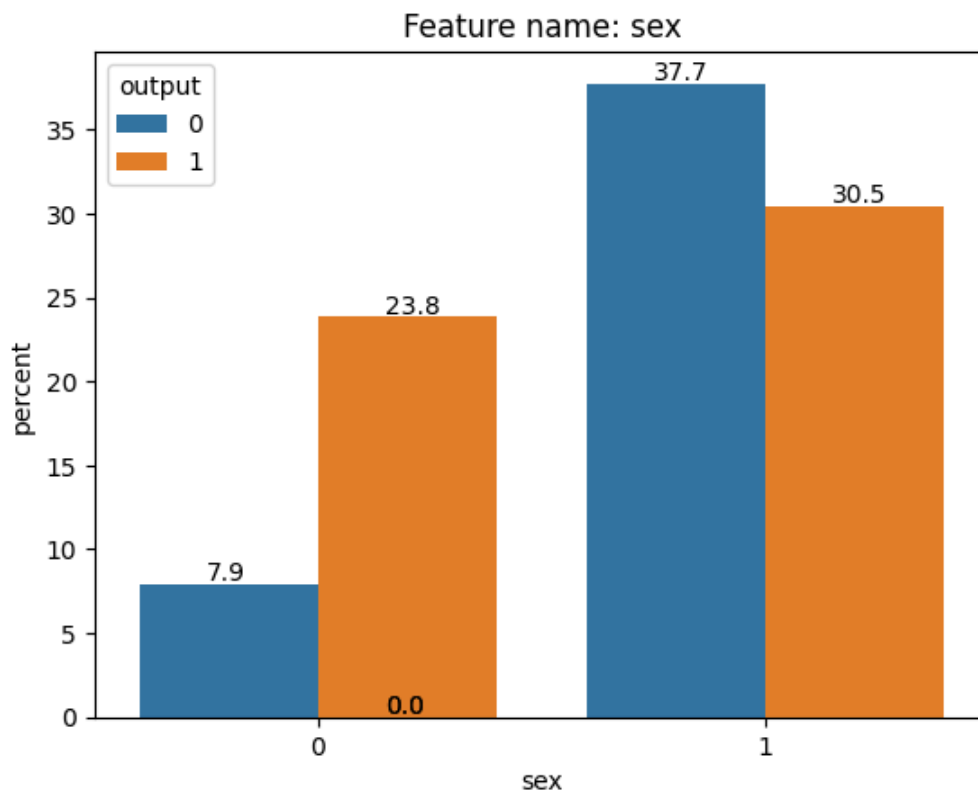
```
<class 'pandas.core.frame.DataFrame'>
Index: 302 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         302 non-null    int64
1   sex         302 non-null    category
2   cp          302 non-null    category
3   trtbps      302 non-null    int64
4   chol        302 non-null    int64
5   fbs         302 non-null    category
6   restecg     302 non-null    category
7   thalachh    302 non-null    int64
8   exng        302 non-null    category
9   oldpeak     302 non-null    float64
10  slp         302 non-null    category
11  caa         302 non-null    int64
12  thall       302 non-null    category
13  output      302 non-null    category
dtypes: category(8), float64(1), int64(5)
memory usage: 20.0 KB
```


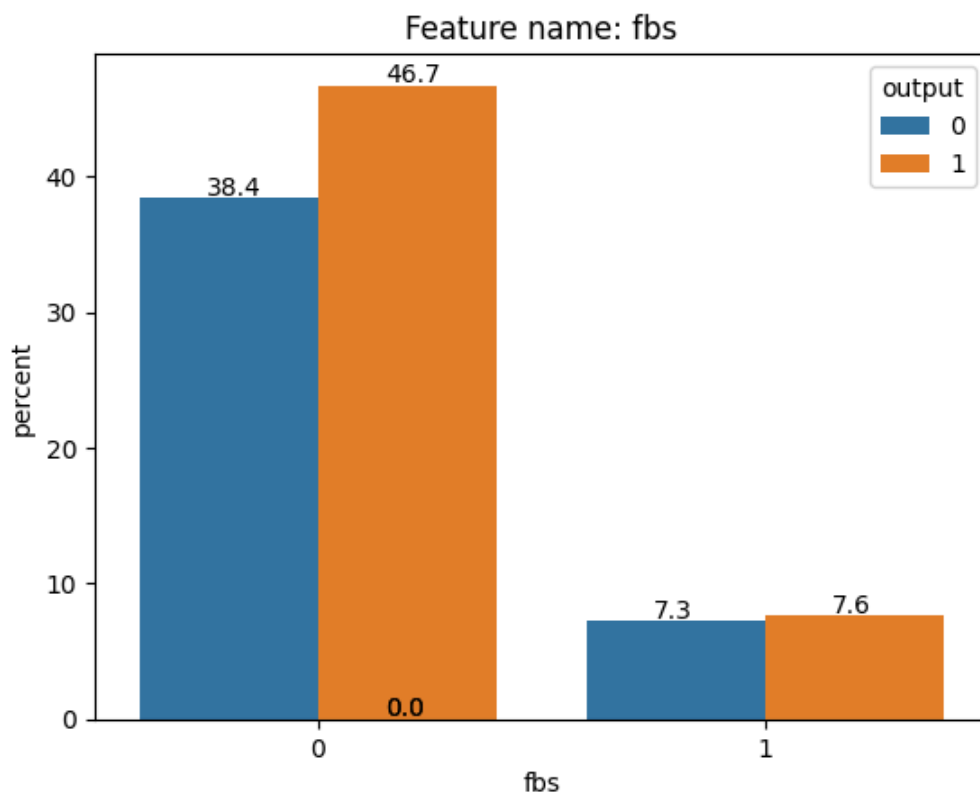
Binary features

The binary features plotted in the count plots below. All them are reasonable balanced. People with high fasting blood sugar (fbs=1) is on the lower side, 15%, but not extremely lopsided that it would need any further action.

```
In [427... bin_features = ['sex', 'fbs', 'exng']

for ftr in bin_features:
    ax = sns.countplot(df_cat, x=ftr, stat="percent", hue='output')
    ### Source for next two lines: https://www.tutorialspoint.com/matplotlib-how-to-show-the-count-values-on-
    for p in ax.patches:
        ax.annotate('{:.1f}'.format(p.get_height()), (p.get_x()+0.15, p.get_height()+0.2))
    ###
    plt.title('Feature name: '+ftr)
    plt.show()
```







Categorical features



The categorical features plotted in the count plots below. Resting ecg ('restecg') and number of major vessels ('caa')

```
In [428... cat_features = ['cp', 'restecg', 'slp', 'thall']

for ftr in cat_features:
    ax = sns.countplot(df_cat, x=ftr, hue='output')
    ### Source for next two lines: https://www.tutorialspoint.com/matplotlib-how-to-show-the-count-values-on-
    for p in ax.patches:
```

```
ax.annotate('{:.1f}'.format(p.get_height()), (p.get_x()+0.05, p.get_height()+0.2))
###
plt.title('Feature name: '+ftr)
plt.show()
```






Real-valued features


```
In [429... real_features = ['age', 'chol', 'trtbps', 'thalachh', 'caa', 'oldpeak']

for ftr in real_features:
    sns.histplot(data=df_cat, x=ftr, hue='output', multiple="dodge")
    plt.title('Feature name: '+ftr)
    plt.show()
```


Feature name: age




Feature name: chol


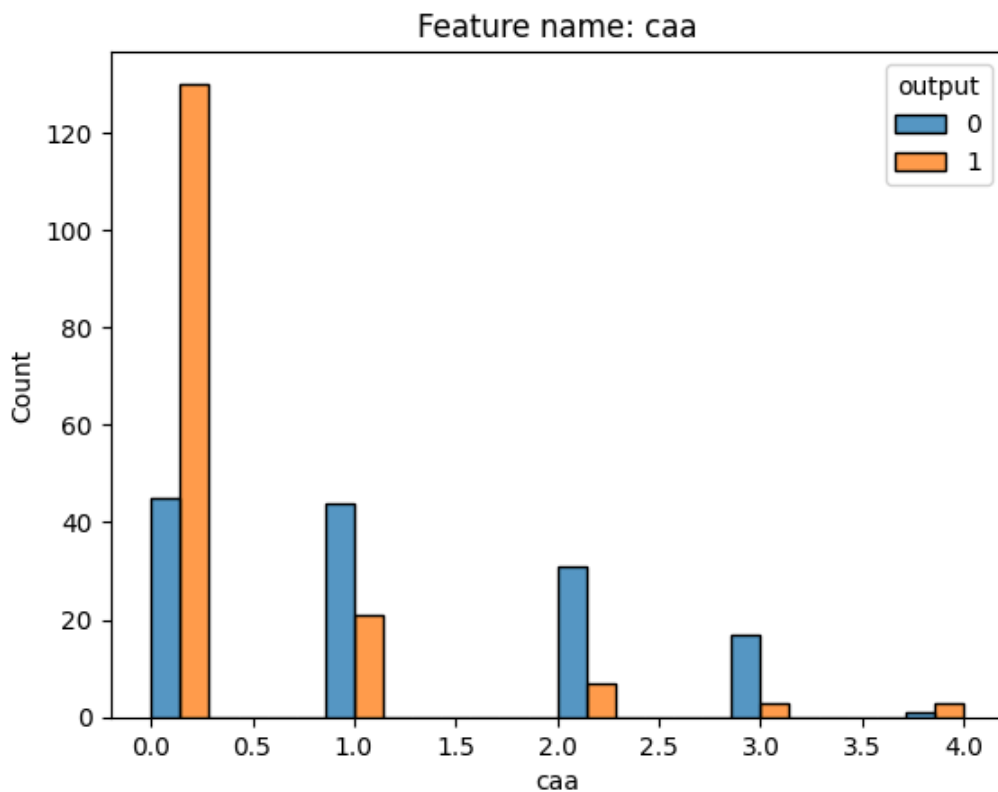


Feature name: trtbps



Feature name: thalachh





The cholesterol feature, 'chol', has some fairly high values that do not fit otherwise gaussian looking distribution. We decided to remove those five rows.

Another issue is that the 'oldpeak' feature in the histogram above looks suspicious. A closer look reveals that it has a lot of zero values for a real-valued feature. Perhaps there's no available data about the previous peak measurement and it was marked as zero.


Note that, in a real world case, the both these actions should be checked with a domain expert who understands the methodology of measuring and interpreting this kind of data.

```
In [430... print('Number of outliers in cholesterol column:', sum(df_cat['chol']>380))

sns.histplot(data=df_cat, x='chol')
df_cat = df_cat.loc[df_cat['chol']<380]
sns.histplot(data=df_cat, x='chol')
plt.title('Feature chol, removing outliers')
plt.legend(['Before cleaning', 'After cleaning'])
plt.show()

print('Before: Number of zero values in "oldpeak" column', sum(df_cat['oldpeak']==0), 'out of', len(df_cat['o
```

Number of outliers in cholesterol column: 5




Before: Number of zero values in "oldpeak" column 98 out of 297

One option is to drop the feature out from the model. Another option is to impute the missing values, which is done below by replacing the zero values with the group mean.

```
In [431... print('Before: Number of zero values in "oldpeak" column', sum(df_cat['oldpeak']==0), 'out of', len(df_cat['o
sns.histplot(data=df_cat, x="oldpeak")
# Impute zero values with the mean
oldpeak_mean = np.mean(df_cat.loc[df_cat['oldpeak']!=0, 'oldpeak'])
df_cat.loc[df_cat['oldpeak']==0, 'oldpeak'] = oldpeak_mean
#print(oldpeak_mean, np.mean(df_cat['oldpeak']))
sns.histplot(data=df_cat, x="oldpeak")
plt.title('Feature name: '+ftr+', updated')
plt.legend(['Before imputing', 'After imputing'])
plt.show()
print('After: Number of zero values in "oldpeak" column', sum(df_cat['oldpeak']==0), 'out of', len(df_cat['ol
```

Before: Number of zero values in "oldpeak" column 98 out of 297



After: Number of zero values in "oldpeak" column 0 out of 297

Now that the data set is clean, we will check how the features correlate with each other.


Correlation analysis of the features plotted in the heatmap below shows no significant correlation between the features (darker red or green would mean correlation).

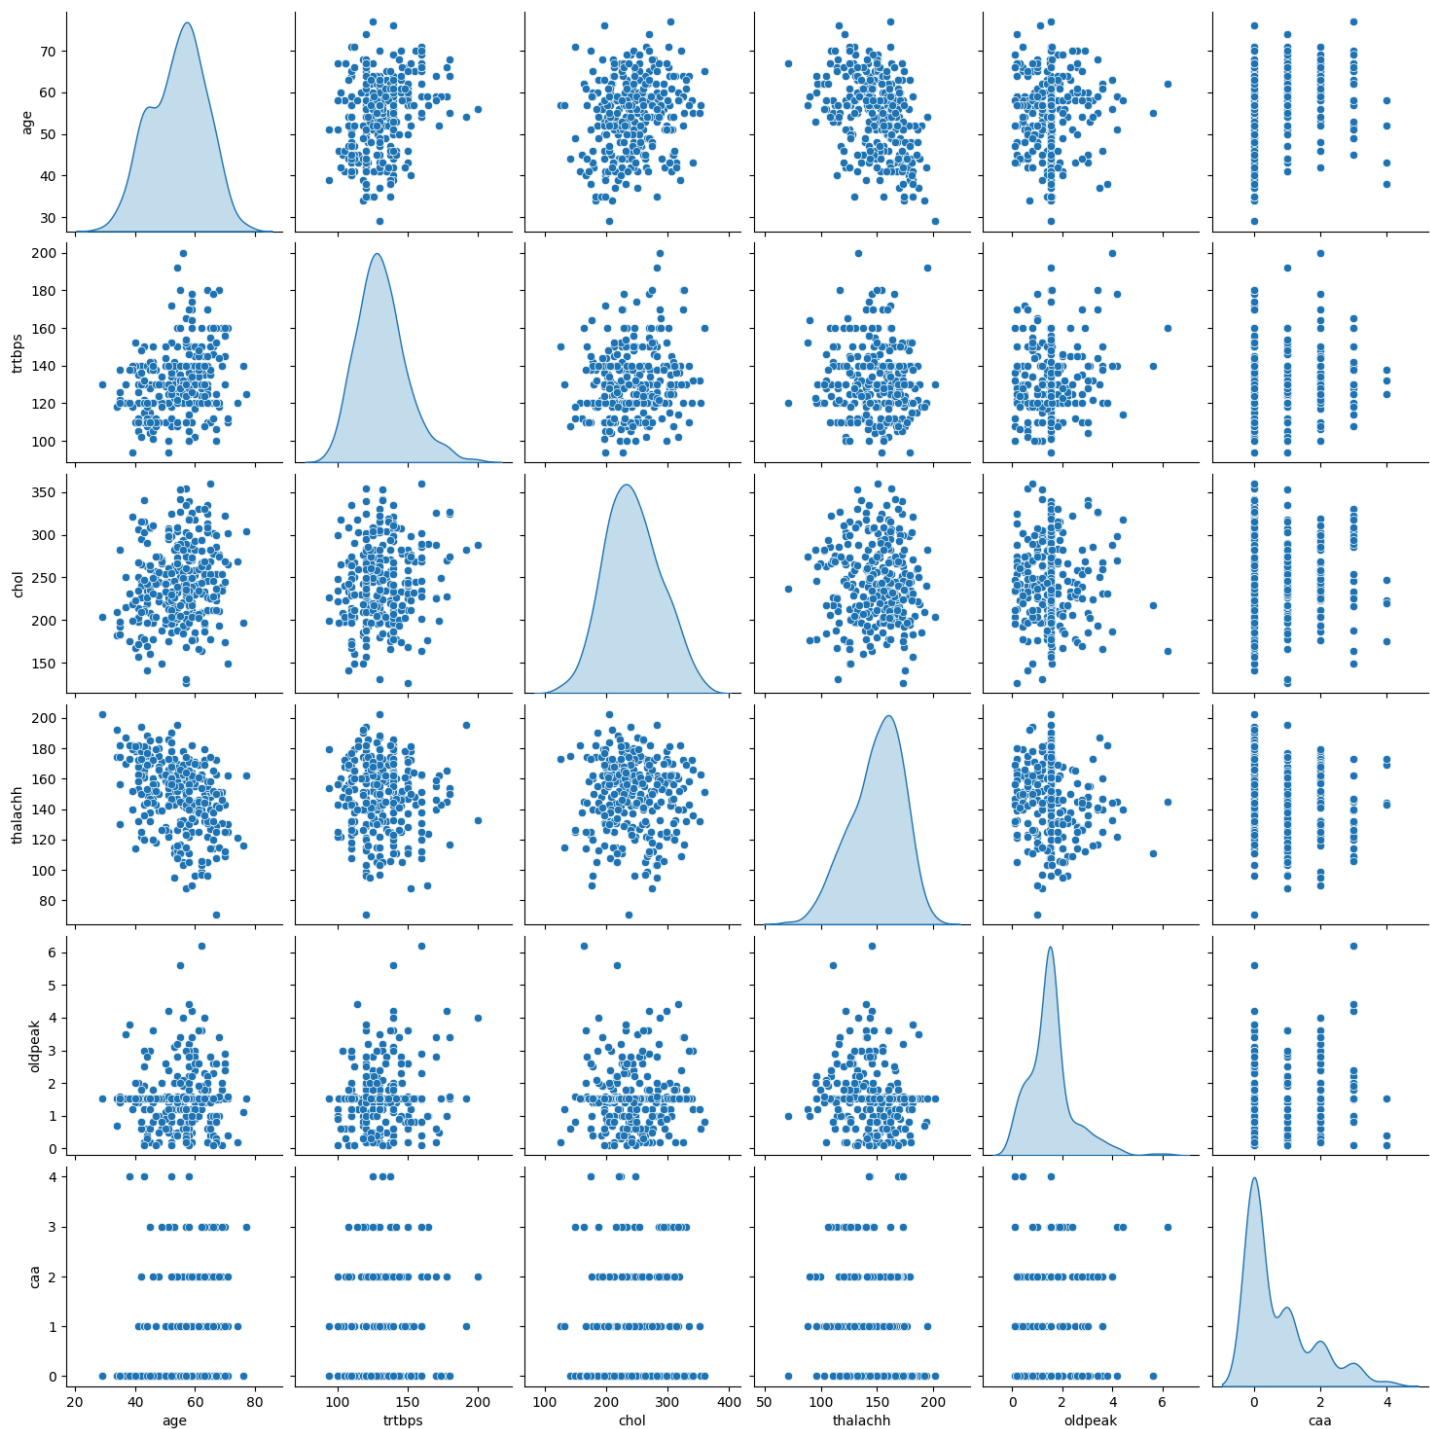
Also the pairplot of the real-valued features does not show significant correlation as the values form roughly round shaped clouds.

In [432...

```
sns.heatmap(df.drop(columns='output').corr(),vmin=-1, vmax=1, cmap='PiYG')
plt.show()

sns.pairplot(df_cat.select_dtypes(['float64', 'int64']), diag_kind='kde')
#sns.pairplot(df_cat.select_dtypes('float64').corr(), diag_kind='kde')
plt.show()
```





Now the data set has been checked for erroneous values and cleaned. We removed a duplicate data row, imputed the zero values in the 'oldpeak' feature and removed a few rows due to outliers in 'chol' feature values.

Prepare training and test sets

Here, we make a 80-20 split for training and test sets.

Also, normalization is performed for the real-valued features: 'age', 'trtbps', 'chol' and 'thalachh'.

```
In [433... from sklearn.model_selection import train_test_split

test_split = 0.2 # 20/80 split test/train

# Normalize continuous variables
columns = ['age', 'trtbps', 'chol', 'thalachh']
for column in columns:
```



```

df[column] = (df[column] - np.mean(df[column])) / np.std(df[column])
df_cat[column] = (df_cat[column] - np.mean(df_cat[column])) / np.std(df_cat[column])

r_state = 6
if False:
    # Original dataframe. Not used.
    df_train, df_test = train_test_split(df, test_size=test_split, random_state=r_state)
elif False:
    # Drop "oldpeak"-feature. Not used.
    df_cat_no_old = df_cat.drop(columns='oldpeak', inplace=False)
    df_train, df_test = train_test_split(df_cat_no_old, test_size=test_split, random_state=r_state)
else:
    # Dataframe with categorical datatypes and normalized continuous variable.
    df_train, df_test = train_test_split(df_cat, test_size=test_split, random_state=r_state)

print('The length of training set:', len(df_train), '\nThe length of testing set:', len(df_test))
y_train = df_train['output']
X_train = df_train.drop(columns='output', inplace=False)
y_test = df_test['output']
X_test = df_test.drop(columns='output', inplace=False)

```

The length of training set: 237
The length of testing set: 60

Models

As mentioned in the intro, four binary classification models are considered,

1. 'Traditional' Logistic Regression,
2. (Pruned) Decision Tree,
3. Ensemble method AdaBoost and
4. Support Vector Machine model. All of these can handle our binary classification output data.

Helper functions

The code below is used for calculating metrics for the models. The metrics extracted are:

- **False Negative Rate**
- **Accuracy**
- F1-score
- Precision
- **Recall**
- Area Under the Receiver Operating Characteristic Curve (ROC AUC)
- **Confusion matrix**

In [434... `from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay, accuracy_score, f1_score, precision_sco`

```

class calc_metrics:
    def __init__(self, y_test, y_pred):
        self.confusion_matrix = None
        self.recall = None
        self.f1_score = None
        self.precision = None
        self.recall = None
        self.FNR = None

        self.get_metrics(y_test, y_pred)

    def get_metrics(self, y_test, y_pred):
        conf_mat = confusion_matrix(y_test, y_pred, labels=[0,1])
        self.confusion_matrix = pd.DataFrame(conf_mat)

```

```

self.accuracy = accuracy_score(y_test, y_pred)
self.f1_score = f1_score(y_test, y_pred)
self.precision = precision_score(y_test, y_pred)
self.recall = recall_score(y_test, y_pred)
self.ROC_AUC = roc_auc_score(y_test, y_pred)
FN = conf_mat[1][0]
P = np.sum(conf_mat[1])
self.FNR = FN/P
print('False Negative Rate:', round(self.FNR,3))
print('Accuracy:', round(self.accuracy, 3))
print('F1-score:', round(self.f1_score,3))
print('Precision:', round(self.precision,3))
print('Recall:', round(self.recall,3))
print('Area Under the Receiver Operating Characteristic Curve (ROC AUC):', round(self.ROC_AUC,3))
print('Confusion matrix:')
disp = ConfusionMatrixDisplay(conf_mat)
disp.plot()
plt.show()

```

Logistic regression model

Logistic regression model created below. Penalty function is L2 and solver is lbfgs. L1 penalty, and liblinear solver were also tried but they yielded slightly worse results and thus are not shown here.

The quality of a logistic regression model suffers when there is interaction between the features. As noted before there were no significant interaction in this data set. If there were, some of the features should be dropped out or dimension reduction tools should be used.

The performance is evaluated with the test set by predicting the outcome and calculating the metrics.

In [435...

```

from sklearn.linear_model import LogisticRegression
LogReg_grid_paras = {'C':np.logspace(-5, 5, num=11, base=2)}
myLogReg = LogisticRegression(penalty='l2', class_weight='balanced', C=0.1, solver='lbfgs', fit_intercept=True)
myLogReg_grid = GridSearchCV(estimator=myLogReg, param_grid=LogReg_grid_paras, scoring=None, n_jobs=None, ref
myLogReg_grid.fit(X_train, y_train) # Create the model
print('Best parameters:', myLogReg_grid.best_params_)
print('Validation accuracy:', myLogReg_grid.best_score_)


y_hat_LogReg = myLogReg_grid.best_estimator_.predict(X_test) # Make prediction
LogReg_metrics = calc_metrics(y_test, y_hat_LogReg) # calculate metrics
df_logreg = pd.DataFrame([X_train.columns, np.round(myLogReg_grid.best_estimator_.coef_,5)[0]], index=['Featu
print(df_logreg.sort_values('Coefficient', axis=1).transpose())

```

```

Best parameters: {'C': np.float64(0.03125)}
Validation accuracy: 0.8187943262411348
False Negative Rate: 0.069
Accuracy: 0.85
F1-score: 0.857
Precision: 0.794
Recall: 0.931
Area Under the Receiver Operating Characteristic Curve (ROC AUC): 0.853
Confusion matrix:

```



	Feature	Coefficient
11	caa	-0.36059
12	thall	-0.32731
1	sex	-0.27261
8	exng	-0.26557
9	oldpeak	-0.25518
4	chol	-0.19611
0	age	-0.113
3	trtbps	-0.08792
5	fbs	-0.03383
6	restecg	0.11064
10	slp	0.3024
7	thalachh	0.38607
2	cp	0.41604

None of the model's coefficients are particularly small. That indicates that most of them make meaningful contribution to the model. The fasting blood sugar feature ('fbs') is only -0.033 making it the least important feature (i.e., closest to zero). Chest pain type ('cp'), maximum heart rate ('thalachh') and number of major vessels ('caa') are strongest predictors.

Decision Tree

Decision tree hyperparameters in the grid search for the best model:

- Max depth, sweep values from 1 to 9.
- CCP alpha, sweep values from 0 to 8.
- Minimum number of samples per leaf, sweep values from 1 to 8.

Total number of models to consider is $9 \times 9 \times 8 = 648$, from which the model with best accuracy is selected.

Performance is evaluated with the test set by predicting the outcome and calculating the metrics.

```
In [436... from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import cross_val_score, GridSearchCV
#from sklearn.ensemble import RandomForestClassifier

DT_grid_paras = {'max_depth':range(1,10), 'ccp_alpha':range(0, 9), 'min_samples_leaf':range(1,9)}
myDT_grid = GridSearchCV(estimator=DecisionTreeClassifier(), param_grid=DT_grid_paras, scoring=None, n_jobs=N
```


```

myDT_grid = myDT_grid.fit(X_train, y_train) # Create the model
print('Best parameters:', myDT_grid.best_params_)
print('Validation accuracy:', myDT_grid.best_score_)

myDT = myDT_grid.best_estimator_ # Save the best model
myDT.fit(X_train, y_train)
y_hat_DT = myDT.predict(X_test) # Estimate using the best model
DT_metrics = calc_metrics(y_test, y_hat_DT) # calculate metrics

```

Best parameters: {'ccp_alpha': 0, 'max_depth': 6, 'min_samples_leaf': 5}
 Validation accuracy: 0.7890070921985816
 False Negative Rate: 0.172
 Accuracy: 0.8
 F1-score: 0.8
 Precision: 0.774
 Recall: 0.828
 Area Under the Receiver Operating Characteristic Curve (ROC AUC): 0.801
 Confusion matrix:



In [437...]

```

df_DT = pd.DataFrame([X_train.columns, np.round(myDT.feature_importances_,5)], index=['Feature', 'Gini import
print(df_DT.sort_values('Gini importance', axis=1).transpose())

```

	Feature	Gini importance
3	trtbps	0.0
5	fbs	0.0
6	restecg	0.0
1	sex	0.00873
9	oldpeak	0.02927
8	exng	0.04175
4	chol	0.04239
10	slp	0.05188
0	age	0.07411
7	thalachh	0.09897
12	thall	0.13303
11	caa	0.15307
2	cp	0.3668

Similarly as with logistic regression Chest pain type ('cp') and number of major vessels ('caa') are strongest predictors.
 Gender ('sex') and fasting blood sugar ('fbs') do not contribute to this model.

AdaBoost

Since AdaBoost algorithm was already coded from scratch as a part of the homework, AdaBoostClassifier from sklearn is used here.

Number of estimators set to 200 after trying a few values.

Performance is evaluated with the test set by predicting the outcome and calculating the metrics.

```
In [438... from sklearn.ensemble import AdaBoostClassifier

myAdaBoost = AdaBoostClassifier(estimator=None, n_estimators=200, learning_rate=0.1, algorithm='SAMME', random_state=42)
myAdaBoost.fit(X_train, y_train) # Create the model

y_hat_AdaBoost = myAdaBoost.predict(X_test) # Estimate using the best model
AdaBoost_metrics = calc_metrics(y_test, y_hat_AdaBoost) # calculate metrics
```

False Negative Rate: 0.069

Accuracy: 0.833


F1-score: 0.844

Precision: 0.771

Recall: 0.931

Area Under the Receiver Operating Characteristic Curve (ROC AUC): 0.836

Confusion matrix:



```
In [439... df_Ada = pd.DataFrame([X_train.columns, np.round(myAdaBoost.feature_importances_,5)], index=['Feature', 'Gini importance'])
print(df_Ada.sort_values('Gini importance', axis=1).transpose())
```

	Feature	Gini importance
5	fbs	0.0
6	restecg	0.0
0	age	0.01795
8	exng	0.0326
1	sex	0.0506
3	trtbps	0.05709
10	slp	0.07118
7	thalachh	0.10444
11	caa	0.12408
12	thall	0.12445
2	cp	0.12609
4	chol	0.13951
9	oldpeak	0.15202

Similarly as with logistic regression and Decision Tree, Chest pain type ('cp') and number of major vessels ('caa') rank high. However, previous peak ('oldpeak') and cholesterol ('chol') are the strongest for AdaBoost. Fasting blood sugar ('fbs') and resting ECG ('restecg') do not contribute to this model.

SVM

Support Vector Machine -model. Trying with two kernels: 'rbf', the radial basis function and 'linear', linear hyperplane function.

Performance is evaluated with the test set by predicting the outcome and calculating the metrics.

Radial Basis Function

Hyperparameters C and gamma are swept using grid search when using the radial basis function. C and gamma are swept independently from 2^{-5} to 2^5 .

Since the performance is fairly poor, linear kernel function is tried next.

```
In [440... from sklearn.svm import SVC

print('\nTry Radial Basis Function kernel:')
SVC_grid_paras = {'C':np.logspace(-5, 5, num=11, base=2), 'gamma':np.logspace(-5, 5, num=11, base=2)}
my_SVC_grid = GridSearchCV(estimator=SVC(), param_grid=SVC_grid_paras, scoring=None, n_jobs=None, refit=True,
my_SVC_grid = my_SVC_grid.fit(X_train, y_train) # Create the model
print('Best parameters:', my_SVC_grid.best_params_)
print('Validation accuracy:', my_SVC_grid.best_score_)
my_SVC = my_SVC_grid.best_estimator_ # Save the best model
y_hat_SVC = my_SVC.predict(X_test)
calc_metrics(y_test, y_hat_SVC) # calculate metrics
```

Try Radial Basis Function kernel:

Best parameters: {'C': np.float64(1.0), 'gamma': np.float64(0.03125)}

Validation accuracy: 0.8315602836879433

False Negative Rate: 0.069

Accuracy: 0.833


F1-score: 0.844

Precision: 0.771

Recall: 0.931

Area Under the Receiver Operating Characteristic Curve (ROC AUC): 0.836

Confusion matrix:



Out[440]... <__main__.calc_metrics at 0x21701d11040>

Linear Function

Only hyperparameter C is swept from 2^{-5} to 2^5 with linear kernel function (does not have gamma parameter).

```
In [441]... SVC_grid_paras = {'C':np.logspace(-5, 5, num=11, base=2)}
my_linSVC_grid = GridSearchCV(estimator=SVC(kernel='linear'), param_grid=SVC_grid_paras, scoring=None, n_jobs=
my_linSVC_grid = my_linSVC_grid.fit(X_train, y_train) # Create the model
print('Best parameters:', my_linSVC_grid.best_params_)
print('Validation accuracy:', my_linSVC_grid.best_score_)
y_hat_linSVC = my_linSVC_grid.best_estimator_.predict(X_test) # Estimate using the best model
SVC_metrics = calc_metrics(y_test, y_hat_linSVC) # calculate metrics
```

Best parameters: {'C': np.float64(0.0625)}

Validation accuracy: 0.8313829787234044

False Negative Rate: 0.069

Accuracy: 0.85


F1-score: 0.857

Precision: 0.794

Recall: 0.931

Area Under the Receiver Operating Characteristic Curve (ROC AUC): 0.853

Confusion matrix:



```
In [442... df_SVC = pd.DataFrame([X_train.columns, np.round(my_linSVC_grid.best_estimator_.coef_[0],5)], index=['Feature
print(df_SVC.sort_values('Coefficient', axis=1).transpose())
```

	Feature	Coefficient
12	thall	-0.51566
8	exng	-0.41879
11	caa	-0.40295
1	sex	-0.37836
4	chol	-0.24572
9	oldpeak	-0.22194
3	trtbps	-0.12261
5	fbs	-0.05791
0	age	-0.0084
6	restecg	0.06408
7	thalachh	0.3504
10	slp	0.45833
2	cp	0.48137

Since we ended up using the linear kernel, we can compare feature importance by looking at the absolute value of the coefficient. The higher the value, the more important the feature.

Similarly as with logistic regression, Decision Tree and AdaBoost, chest pain type ('cp') and number of major vessels ('caa') rank high. Thal rate ('thall') and slope ('slp') are also strongest here. Fasting blood sugar ('fbs'), age and resting ECG ('restecg') have minimal effect in this model.

Results and Analysis

```
In [443... print('Logistic Regression accuracy:',round(LogReg_metrics.accuracy,3))
print('Decision Tree accuracy:',round(DT_metrics.accuracy,3))
print('AdaBoost accuracy:',round(AdaBoost_metrics.accuracy,3))
print('Support Vector Machine accuracy:',round(SVC_metrics.accuracy,3))
print()
print('Logistic Regression recall:',round(LogReg_metrics.recall,3))
print('Decision Tree recall:',round(DT_metrics.recall,3))
print('AdaBoost recall:',round(AdaBoost_metrics.recall,3))
print('Support Vector Machine recall:',round(SVC_metrics.recall,3))
```



```

#LogReg_metrics.accuracy
df_results = pd.DataFrame(np.round([[LogReg_metrics.accuracy, DT_metrics.accuracy, AdaBoost_metrics.accuracy,
[LogReg_metrics.precision, DT_metrics.precision, AdaBoost_metrics.precision, SVC_metrics.precis
[LogReg_metrics.f1_score, DT_metrics.f1_score, AdaBoost_metrics.f1_score, SVC_metrics.f1_score]
[LogReg_metrics.recall, DT_metrics.recall, AdaBoost_metrics.recall, SVC_metrics.recall],
[LogReg_metrics.ROC_AUC, DT_metrics.ROC_AUC, AdaBoost_metrics.ROC_AUC, SVC_metrics.ROC_AUC]],3)
index=['accuracy', 'precision', 'f1_score', 'recall', 'ROC AUC'], columns=['Logistic Regression

df_FNR = pd.DataFrame(np.round([[LogReg_metrics.FNR, DT_metrics.FNR, AdaBoost_metrics.FNR, SVC_metrics.FNR]],
index = ['FNR'], columns=['Logistic Regression', 'Decision Tree', 'AdaBoost', 'SVM'])


#print(df_FNR)
sns.barplot(df_FNR)
plt.title('False Negative Rate')
plt.show()


plt.plot(df_results.transpose())
plt.legend(df_results.index)
plt.show()

```

Logistic Regression accuracy: 0.85
 Decision Tree accuracy: 0.8
 AdaBoost accuracy: 0.833
 Support Vector Machine accuracy: 0.85

Logistic Regression recall: 0.931
 Decision Tree recall: 0.828
 AdaBoost recall: 0.931
 Support Vector Machine recall: 0.931





Discussion and conclusion

When evaluating the model options, one should pay close attention to false negatives (bottom left corner of the confusion matrix). Considering the type of data where we want to predict heart conditions, we do not want to miss any potentially positive cases. That is why close attention should be paid to getting a low False Negative Rate (FNR) or, equivalently, high Recall score ($1 - \text{FNR}$).

Overall, the differences between the methods are relatively small and small changes in train-test split or random seed can change the models' ranking. The fact that logistic regression performs pretty well and SVM with linear kernel is at least as good as with non-linear kernel suggest that most of the features have close to linear relationship to the outcome.

Based on the metrics logistic regression and Support Vector Machine with linear kernel perform the best and AdaBoost is also strong. I would recommend using the logistic regression for this data because it has superior explainability.

Training time

Approximate training times are listed below. Decision tree has three-hyperparameter-grid and SVM with RBF has two, making them much slower.

- Logistic Regression: <2s
- Decision Tree: 30s
- AdaBoost: <3s
- SVM:
 - RBF: 10s
 - Linear: 1s

The data set was relatively small so training time for all the algorithms was reasonable.